

CS598 : HW1_apoorva6

Apoorva

9/8/2019

Contents

Question 1 - KNN	1
Question 2 - Linear Regression through Optimization	4

Question 1 - KNN

Write an R function to fit a KNN regression model. Complete the following steps

Part a.

Write a function `myknn(xtest, xtrain, ytrain, k)` that fits a KNN model that predict a target point or multiple target points `xtest`. Here `xtrain` is the training dataset covariate value, `ytrain` is the training data outcome, and `k` is the number of nearest neighbors. Use the '2 norm to evaluate the distance between two points. Please note that you cannot use any additional R package within this function.

#Function to find Euclidean distance between test and train variables

```
euclideanDist <- function(a, b){
  eu_dist = 0
  for(k in c(1:(length(a))))
  {
    eu_dist = eu_dist + (a[[k]]-b[[k]])^2
  }
  eu_dist = sqrt(eu_dist)
  return(eu_dist)
}

# My Knn function
myknn <- function(xtest, xtrain, ytrain, k)
{
  pred =c()

  for(i in c(1:nrow(xtest)))
  {
    dist = c()
    for(j in c(1:nrow(xtrain)))
    {
      dist <- c(dist, euclideanDist(xtest[i,], xtrain[j,]))
    }
    dist_df <- data.frame(dist = dist,y = ytrain)
    dist_eu <- dist_df[order(dist_df$dist),]
    dist_eu_k <- dist_eu[1:k,]
    avg_dist <- mean(dist_eu_k[, "y"])
    pred <- c(pred, avg_dist)
  }
}
```

```
mean_Sq_error <- mean((pred - test[,6])^2)
return(mean_Sq_error)
}
```

Part b.

Generate 1000 observations from a five-dimensional normally distribution:

$$\mathcal{N}(\mu, \Sigma_{5 \times 5})$$

where $\mu = (1, 2, 3, 4, 5)^T$ and $\Sigma_{5 \times 5}$ is an autoregressive covariance matrix, with the (i, j) th entry equal to $0.5^{|i-j|}$. Then, generate outcome values Y based on the linear model

$$Y = X_1 + X_2 + (X_3 - 2.5)^2 + \epsilon$$

where ϵ follows i.i.d. standard normal distribution. Use `set.seed(1)` right before you generate this entire data. Print the first 3 entries of your data.

```
library(MASS)
set.seed(1)
# Mu
mu <- c(1,2,3,4,5)
# Creating the matrix for sigma
calc2 <- vector()
for (i in 1:5)
{
  for (j in 1:5)
  {
    calc <- 0.5 ^ abs(i-j)
    calc2 <- c(calc2, calc)
  }
}
Sigma <- matrix(calc2,5,5)

X <- mvrnorm(n = 1000, mu, Sigma)
colnames(X) <- c("X1", "X2", "X3", "X4", "X5")

# Generating outcome values 'Y'
eps <- rnorm(1000, mean=0, sd=1)
Y <- X[, "X1"] + X[, "X2"] + ((X[, "X3"] - 2.5)^2) + eps

data <- cbind(X, Y)
data[1:3,]
```

```
##           X1           X2           X3           X4           X5           Y
## [1,]  2.0770490 3.555163 2.641969 3.902436 5.108741 4.135994
## [2,]  2.4780195 2.161175 2.188487 2.796376 5.330744 5.365376
## [3,] -0.1413538 2.630428 4.666608 4.493909 5.698190 5.505069
```

Part c.

Use the first 400 observations of your data as the training data and the rest as testing data. Predict the Y values using your KNN function with $k = 5$. Evaluate the prediction accuracy using mean squared error

$$\frac{1}{N} \sum_i (y_i - \hat{y}_i)^2$$

```
train <- data[1:400,]
#nrow(train)
test <- data[401:1000,]
#nrow(test)

# The function Myknn returns MSE value
myknn(test[, -6], train[, -6], train[, 6], 5)

## [1] 2.191387
```

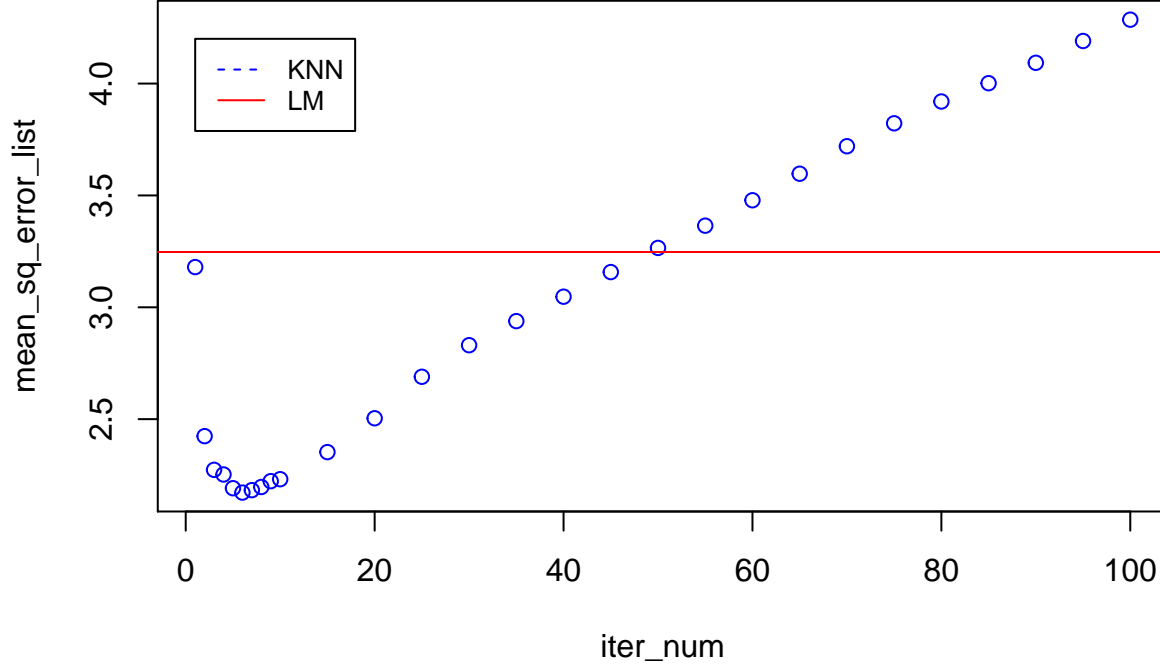
Part d.

Compare the prediction error of a linear model with your KNN model. Consider k being 1, 2, 3, . . . , 9, 10, 15, 20, . . . , 95, 100. Demonstrate all results in a single, easily interpretable figure with proper legends.

```
#Regression function
reg <- lm(Y~., data=as.data.frame(train))
prediction_lm <- predict(reg, as.data.frame(test))
mean_sq_error_lm <- mean((prediction_lm - test[, 6])^2)

# Note down MSE values for values of K ranging from 1 to 100 in steps of 5
mean_sq_error_list <- c()
iter_num <- c(1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 15, 20, 25, 30, 35, 40, 45, 50, 55, 60, 65, 70, 75, 80, 85, 90, 95, 100)
for (i in iter_num)
{
  mse <- myknn(test[, -6], train[, -6], train[, 6], i)
  mean_sq_error_list <- c(mean_sq_error_list, mse)
}

#Plot of MSE of the linear regression model and the knn model for different values of k
plot(mean_sq_error_list~iter_num, col = "blue", xlim = c(1, 100))
abline(h=mean_sq_error_lm, col="red")
legend(1, 4.2, legend = c("KNN", "LM"),
      col= c("blue", "red"), lty=2:1, cex=0.8)
```



Question 2 - Linear Regression through Optimization

Linear regression is most popular statistical model, and the core technique for solving a linear regression is simply inverting a matrix:

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

However, let's consider alternative approaches to solve linear regression through optimization. We use a gradient descent approach. We know that $\hat{\beta}$ can also be expressed as

$$\hat{\beta} = \arg \min \ell(\beta) = \arg \min \frac{1}{2n} \sum_{i=1}^n (y_i - x_i^T \beta)^2.$$

And the gradient can be derived

$$\frac{\partial \ell(\beta)}{\partial \beta} = -\frac{1}{n} \sum_{i=1}^n (y_i - x_i^T \beta) x_i.$$

To perform the optimization, we will first set an initial beta value, say $\beta = \mathbf{0}$ for all entries, then proceed with the updating

$$\beta^{\text{new}} = \beta^{\text{old}} - \frac{\partial \ell(\beta)}{\partial \beta} \times \delta,$$

where δ is some small constant, say 0.1. We will keep updating the beta values by setting β^{new} as the old value and calculating a new one until the difference between β^{new} and β^{old} is less than a prespecified threshold ϵ , e.g., $\epsilon = 10^{-6}$. You should also set a maximum number of iterations to prevent excessively long running time.

We will implement our function on the Boston Housing data from the `mlbench` package. We will remove `medv`, `town` and `tract` from the data and use `cmedv` as the outcome. We will use a scaled and centered version of the data for estimation.

```
library(mlbench)
data(BostonHousing2)
X = BostonHousing2[, !(colnames(BostonHousing2) %in% c("medv", "town", "tract", "cmedv"))]
X = data.matrix(X)
X = scale(X)
Y = as.vector(scale(BostonHousing2$cmedv))
```

Part a.

Based on the above description, write your own R function `mylm_g(x, y, delta, epsilon, maxitr)` to implement this optimization version of linear regression. The output of this function should be a vector of the estimated beta value.

```
mylm_g <- function(x, y, delta, epsilon, maxitr)
{
  iter <- 1
  diff <- 10
  beta_old <- matrix(0, nrow=15)
  while((iter < maxitr) & (diff > epsilon))
  {
    doh_beta <- t(x) %*% ((x) %*% beta_old - y) / length(y)
    beta_new <- beta_old - (doh_beta * delta)
    diff <- sqrt(sum((beta_old-beta_new)^2))
    beta_old <- beta_new
    iter <- iter + 1
  }
  print(beta_old)
}
```

```
my_lm <- mylm_g(X,Y,0.1,(10^(-6)),400)
```

```
##           [,1]
## lon      -0.033218658
## lat       0.029718413
## crim     -0.096323627
## zn        0.115461478
## indus     0.003275932
## chas      0.072240723
## nox      -0.197124711
## rm        0.288555314
## age       0.006776135
## dis      -0.320624437
## rad       0.270922202
## tax      -0.214677567
## ptratio  -0.205533203
## b         0.091161291
## lstat    -0.417368861
```

Part b.

Compare your results with the `lm()` function on the same data. Experiment on different `maxitr` values to obtain a good solution

```
##### Linear regression using library #####
reg <- lm(Y~X)
lm_coef <- summary(reg)$coefficients[2:16,1]

##### Comparision #####
beta_coef <- data.frame(my_lm,lm_coef,my_lm-lm_coef)
colnames(beta_coef) <- c("GradientDescent","Lm","Difference")
beta_coef
```

##	GradientDescent	Lm	Difference
## lon	-0.033218658	-0.032316441	-9.022174e-04
## lat	0.029718413	0.030245087	-5.266747e-04
## crim	-0.096323627	-0.097935969	1.612342e-03
## zn	0.115461478	0.118273098	-2.811619e-03
## indus	0.003275932	0.011390378	-8.114446e-03
## chas	0.072240723	0.071312253	9.284695e-04
## nox	-0.197124711	-0.199703772	2.579061e-03
## rm	0.288555314	0.287232811	1.322503e-03
## age	0.006776135	0.007564852	-7.887173e-04
## dis	-0.320624437	-0.321039342	4.149052e-04
## rad	0.270922202	0.290850755	-1.992855e-02
## tax	-0.214677567	-0.236526155	2.184859e-02
## ptratio	-0.205533203	-0.206804965	1.271762e-03
## b	0.091161291	0.091235409	-7.411765e-05
## lstat	-0.417368861	-0.417972819	6.039576e-04